



Mobile Robot Algorithms Laboratory

Lab 4

Thursday Week 4

<http://www.andrew.cmu.edu/course/16-362-862>



Agenda Today

CarnegieMellon
THE ROBOTICS INSTITUTE

- Administrative Issues
- Feedback Control
- Tuning
- Model Reference Control
- Identification
- Feedforward Control
- 2 Dof Control
- Overview of Lab 4



Agenda Today

Carnegie Mellon
THE ROBOTICS INSTITUTE

- **Administrative Issues**
- Feedback Control
- Tuning
- Model Reference Control
- Identification
- Feedforward Control
- 2 Dof Control
- Overview of Lab 4



Agenda Today


Carnegie Mellon
THE ROBOTICS INSTITUTE

- Administrative Issues
- **Feedback Control**
- Tuning
- Model Reference Control
- Identification
- Feedforward Control
- 2 Dof Control
- Overview of Lab 4



Control Matters

Carnegie Mellon
THE ROBOTICS INSTITUTE





Control Matters

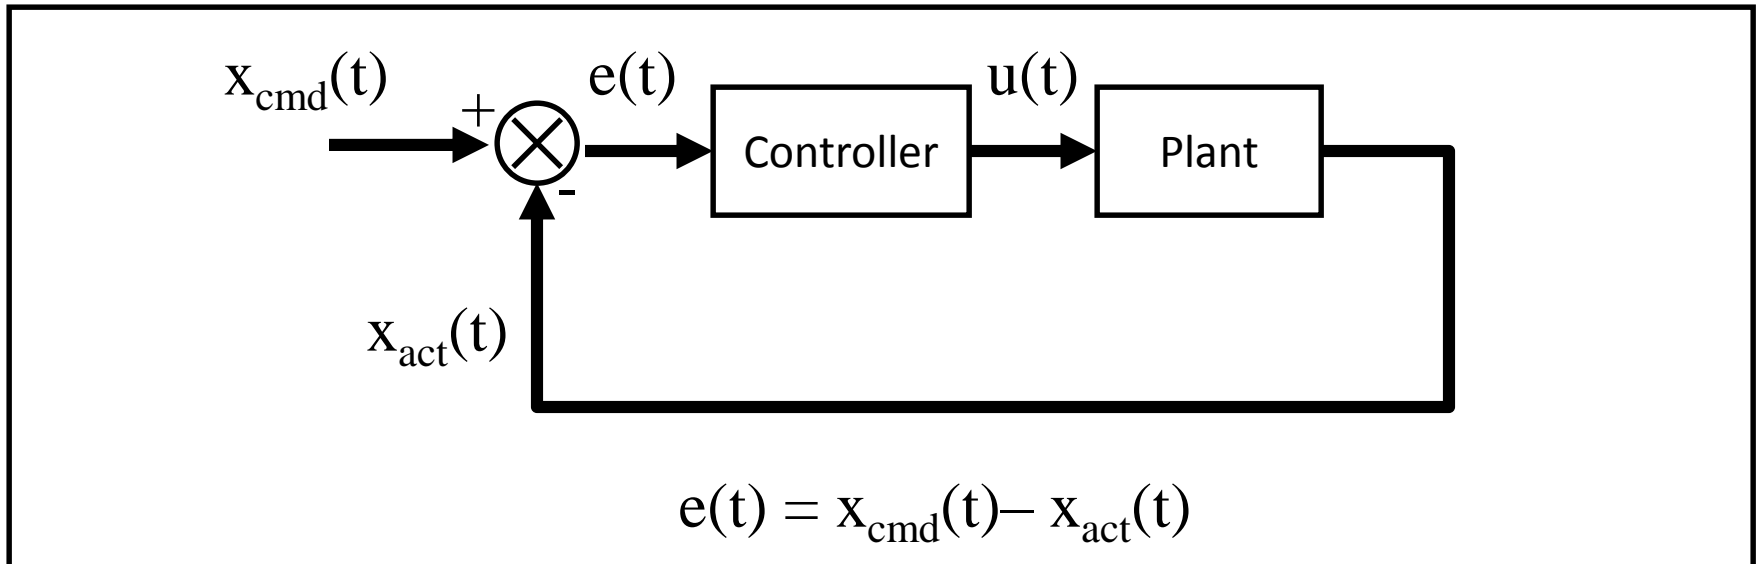
Carnegie Mellon
THE ROBOTICS INSTITUTE





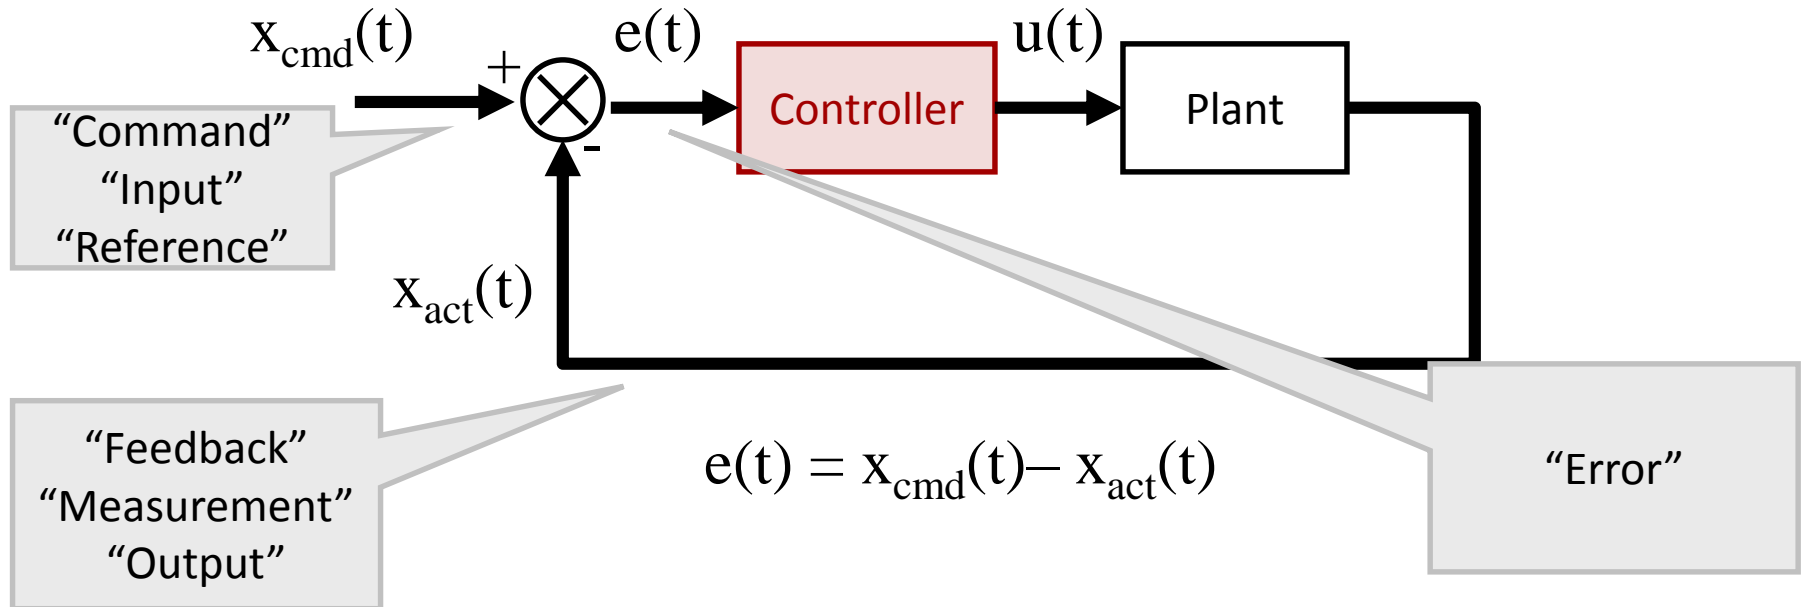
FeedBack Control Primer

- Contrast with “open loop” (send and forget) control.
- Philosophy: You can’t predict the future well in most cases (disturbances, model errors).
- Approach: Actively monitor what is going on and correct for errors that were not or could not be predicted.
- Basic issue: $u(t)$ (control) is not $x(t)$ (state). Often $u(t)$ is related to derivatives of $x(t)$. IOW, system has dynamics.





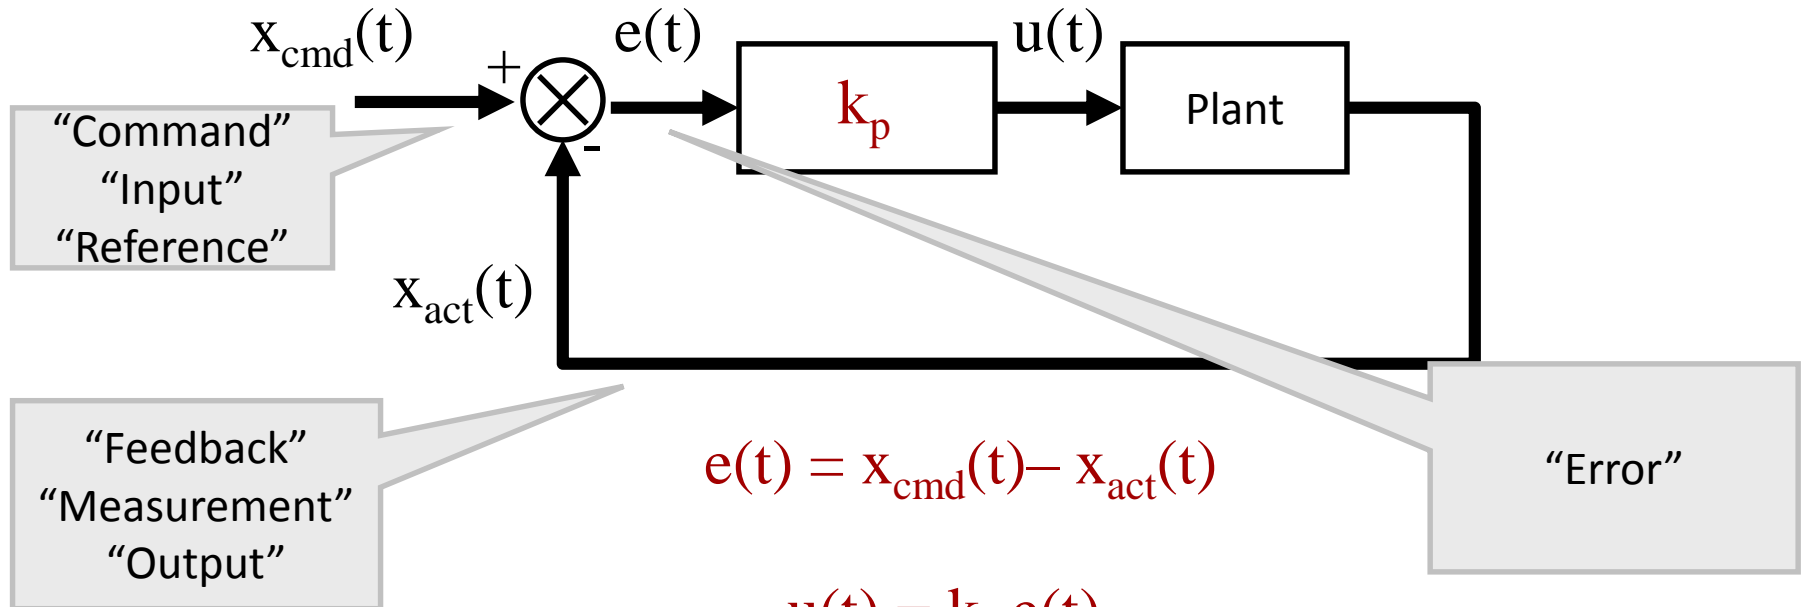
“Controllers”



- So... **What’s in the controller box?**



“Controllers”




$$e(t) = x_{cmd}(t) - x_{act}(t)$$

$$u(t) = k_p e(t)$$

Basic Proportional Control

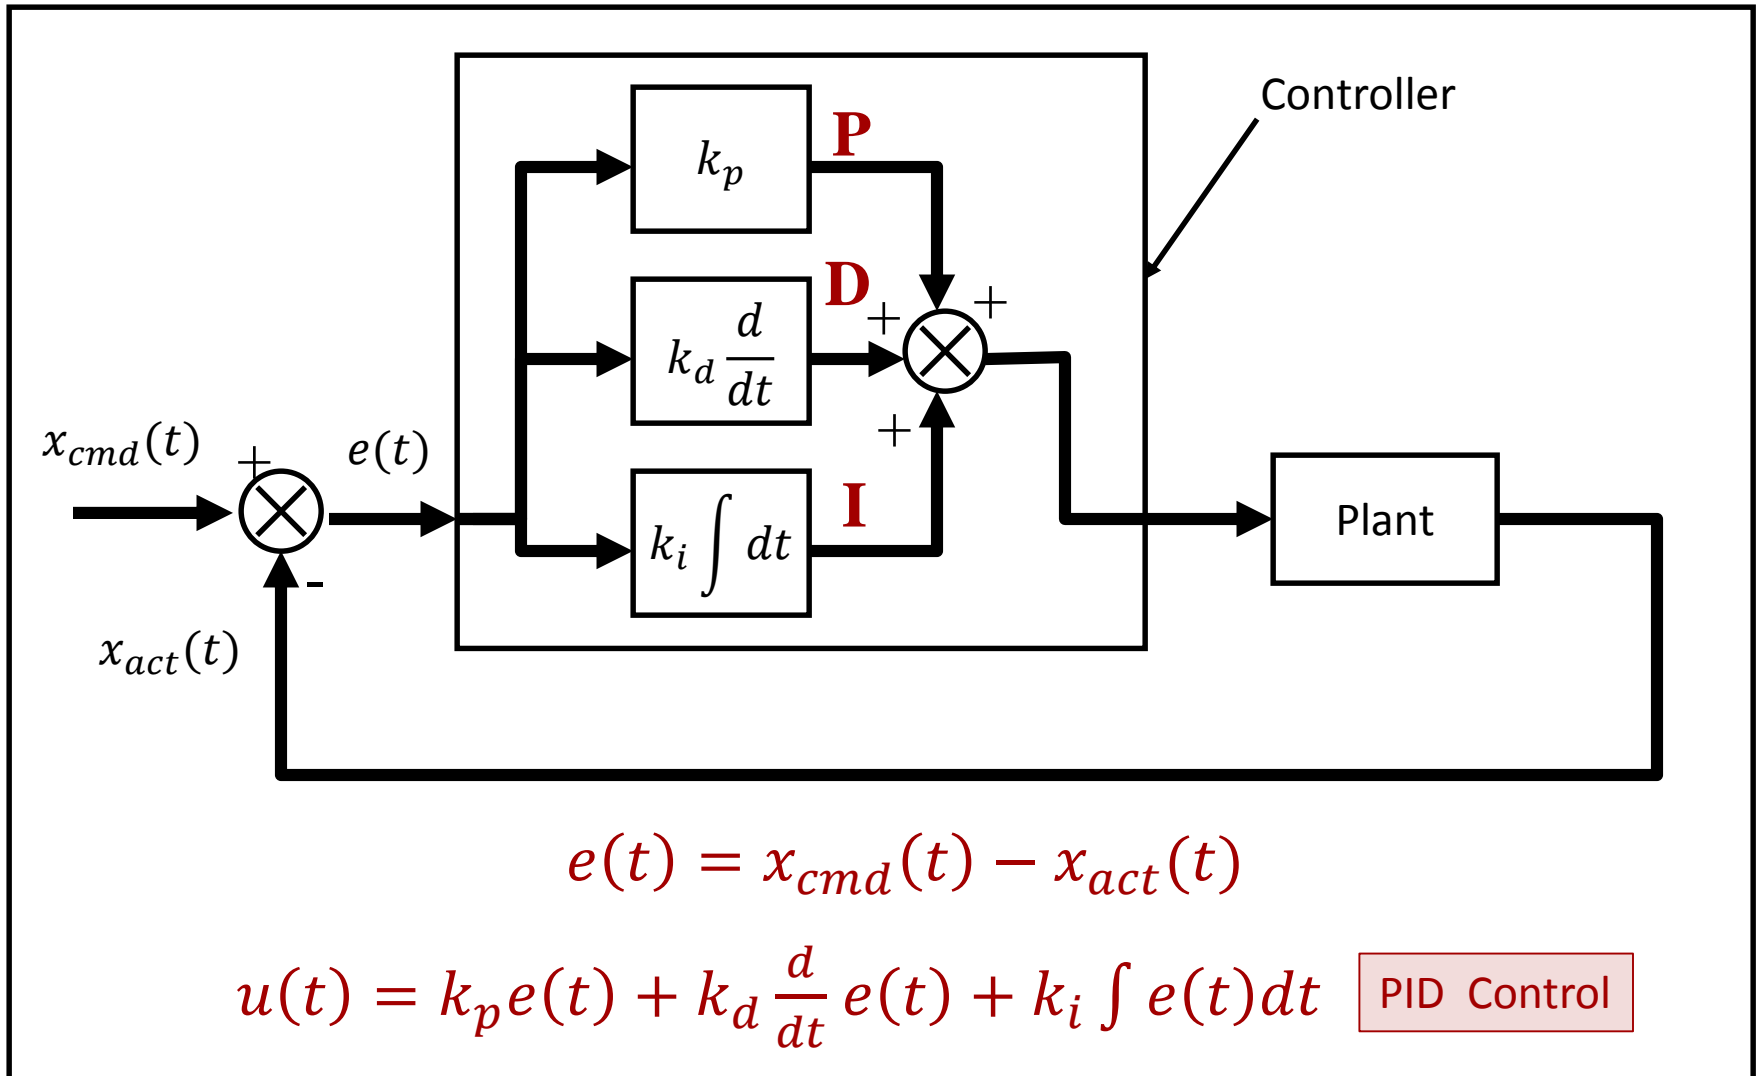
- Generates a “restoring effort”
- “Stabilizes” the system
- like gravity stabilizes a pendulum.





Feedback Control Primer – PID Control



Carnegie Mellon
THE ROBOTICS INSTITUTE





Example

- Feedback Control of Point Mass.





D Term

- Basic idea is to drive the error to zero.
- However:
 - For the same instantaneous value of $e(t)$ shouldn't we **back off** on $u(t)$ **if** the error is **decreasing**?
 - Conversely, shouldn't we **add some more** to $u(t)$ **if** the error is **increasing**?
- Reward the correct **trend** ...
 - reduce overshoot.
 - reduce settling time.

$$u(t) = k_p e(t)$$

$$u(t) += k_d \frac{d}{dt} e(t)$$

a) So decrease the control a little.

b) So increase the control a little

a) < 0 when making Progress.

b) > 0 when not making progress



I Term

- Basic idea is to drive the error to zero.
- However:
 - Shouldn't we **increase** $u(t)$ **if the error has persisted** for a long time.
 - Shouldn't we **decrease** $u(t)$ **if the error has NOT persisted** for a long time.
- Penalize persistence and reduce steady state error.

$$u(t) = k_p e(t)$$

$$u(t) += k_i \int e(t) dt$$

a) So increase the control a lot.

b) So increase control a little

a) large when error persists.

b) small when error is brief



Feedback Control Primer – Implementation In Java

```
// Form error and its integral and derivative
lastError = error; // DECLARE APPROPRIATELY
error      = goalState - actualState;
// COMPUTE DELTIME APPROPRIATELY
errorDerivative = (error - lastError) / delTime;
errorIntegral   += error * delTime;

// Clamp the integral to avoid windup
double sign = errorIntegral > 0.0f ? 1 : -1;
if(Math.abs(errorIntegral) > errorIntegralMax)
    errorIntegral = sign * errorIntegralMax;

// Compute desired control
double control =          error * p_gain
                    + errorDerivative * d_gain
                    + errorIntegral   * i_gain;
```



Feedback Control Primer – Caveats – Measuring dt

- To avoid dependence on processor load, use **actual measurements of $delTime$** to make algorithm work for any $delTime$.
 - Simple approximation is to use `sleep()`.
 - **Q: When does that work?**
- State variable (x) being controlled may be...
 - position, velocity, angle, flow, force, etc.
- Control variable (u) may be ...
 - power, torque, speed, etc.
- Each case for x and u will require different tuning.



Feedback Control Primer –

Caveats - Windup

- Beware integrals in real-time control systems subject to **externally introduced persistent errors**.
 - E.g. computing elapsed time across a breakpoint (implicit integral)
 - E.g. Error integrals in PID control (explicit integral)
- “Windup” is often **VERY DANGEROUS !!!!**
 - Dangerous just like **positive feedback**
 - Get this right the first time
- Make sure to clear all integrals at the appropriate time.
 - Don’t remember integrals from one run to the next.
 - Clear the integral AFTER you press the go button
 - → not when the GUI is constructed.



Uncommanded Motion

- The error, error rate etc. can provide a good basis for detecting unsafe conditions.
- Uncommanded motion means either:
 - A) I **am** moving and I **am not** supposed to be
 - B) I **am not** moving and I **am** supposed to be.
- Generally, check if:

$$|v_{cmd} - v_{act}| > threshold$$



Loss of Control Authority

Carnegie Mellon
THE ROBOTICS INSTITUTE

- If error has changed by a huge amount since last time, something is not right.
- Could be **power**:
 - someone hit the Estop button
 - Generator or engine shut off
- Could be **traction / locomotion**
 - someone picked up the robot
 - wheels are slipping / robot blocked by obstacle
- Could be **error**
 - Cable cut
 - Connector fell off
- You could check $\frac{d^2}{dt^2} e(t)$ to detect these things



Feedback Control Primer – Bad News

- Its not all good news.
- Proportional Term
 - Cannot remove steady state error
 - when there is deadband (e.g. stiction)
 - when there is motion (e.g. tracking)
- Integral Term
 - Introduce to eliminate steady state error.
 - But, causes overshoot & oscillation.
- Derivative term
 - reacts strongly to noise, causes jerky motion.
- Worse yet, the effects are coupled. Cannot tune very incrementally.



Feedback Control Primer – Hints

- PID is but one option. PD is another. P is a third. There are many other control schemes. In rough order of sophistication:
 - Cascade
 - Gain scheduling
 - Reference model based
 - Feedforward
 - Two degree of freedom – Stay Tuned for this
 - Predictive
 - Optimal
 - Nonlinear
 - Adaptive
- For PID.
 - Get P working.
 - Add D to reduce settling time.
 - Add I if steady state error is an issue.
- Every term addition or gain change changes everything.



Feedback Control Primer – Hints

- “error” is not always error. It may be caused by unreasonable expectations.
 - Only force / acceleration can change instantaneously in nature
 - Position and velocity (and hence potential and kinetic energy and momentum) **cannot**.
- Don't **expect** the system to do what it cannot do, errors blow up and chaos results.
 - Think about error in terms of **deviation from reasonable response** motions.
- **See model reference / feedforward control later.**



Agenda Today

Carnegie Mellon
THE ROBOTICS INSTITUTE

- Administrative Issues
- Feedback Control
- **Tuning**
- Model Reference Control
- Identification
- Feedforward Control
- 2 Dof Control
- Overview of Lab 4



Tuning – Best Practices

Carnegie Mellon
THE ROBOTICS INSTITUTE

- Be **structured** and **deliberate**.
- Best gains and thresholds depend on..
 - Momentum, which depends on....
 - Speed of motion, which depends on...
 - The **distance to the goal**
 - So, performance depends on the task
 - The gains themselves a moment ago
 - So, performance now depends on performance earlier
 - Its circular because real systems have dynamics




Tuning - Robustness

- **Know** the limits of your tuning.
 - Sometimes, double the speed and it all falls apart.
- **Reduce** the limits of your tuning.
 - A robust “get me somewhere” is pretty useful.
- Regression testing
 - Keep your test code operational.



Tuning – Empiricism


- Don't try to predict what can be measured.
- Example:
 - Determine error threshold for servo shutoff experimentally
 - disable the shutoff
 - plot the error





Tuning

- Turning 90 deg
- $P=1, I,D=0$
- Pretty good first guess.
- Takes 3-4 seconds to get there.
- Enter **greed**.
 - Can we make it faster?
 - Increase P gain?




↑
3 secs



Tuning

- Turning 90 deg
- $P=2$, $D=0.1$, $I=0$
- 3 times faster.
- But: now it overshoots.
- Try to add some D to reduce overshoot.




1.5 secs



Tuning

- Turning 90 deg
- $P=2$, $D=0.3$, $I=0$
- Too much D leads to slamming on the brakes early.
 - Robot seems to bounce off the goal.




2.0 secs



Tuning

- Turning 90 deg
- $P=2$, $D=0.15$, $I=0$
- Best compromise.



1.25 secs



Tuning

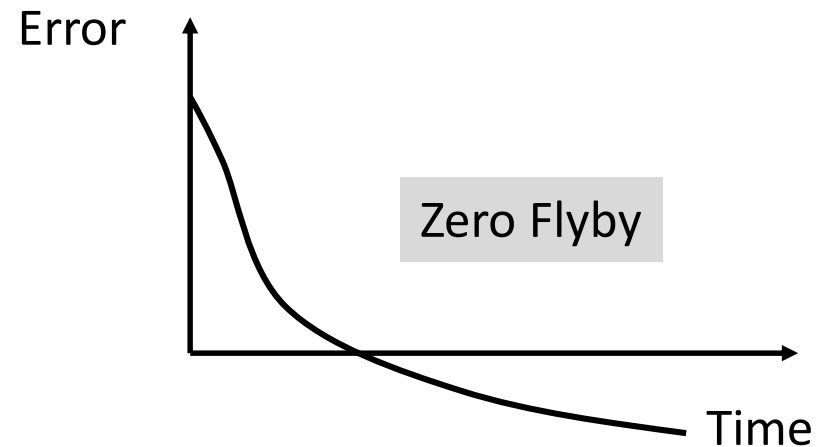

- Its not about beautiful code!
 - Beautiful code helps you maintain it but it won't prop up a fundamentally bad approach.
 - There is no point maintaining it if it does not work.
- Behavior is hard to understand and hard to quantify.
 - Get QUANTITATIVE when possible. It beats subjective assessments in most cases.
- Basic approach is wait for it
 - R.E.C.O.R.D..... S.O.M.E..... D.A.T.A.
 - And study it
- **MATLAB makes this easy**



Termination & Convergence

Carnegie Mellon
THE ROBOTICS INSTITUTE

- In general, you should check error velocity as well as error itself. Why?
 - Could have a lucky sample while flying past the goal
- If you do that, might as well implement the derivative term of PID.
 - Can always set the gain to zero.







Beware Last Cycle Memory

Carnegie Mellon
THE ROBOTICS INSTITUTE



- **Last cycle memory** (for derivatives and integrals) **can be tricky** when starting and stopping servos.
- For integrators, the key thing is that the integrator must be **either or both of**:
 - shut off
 - nulled





Beware Derivatives Too...

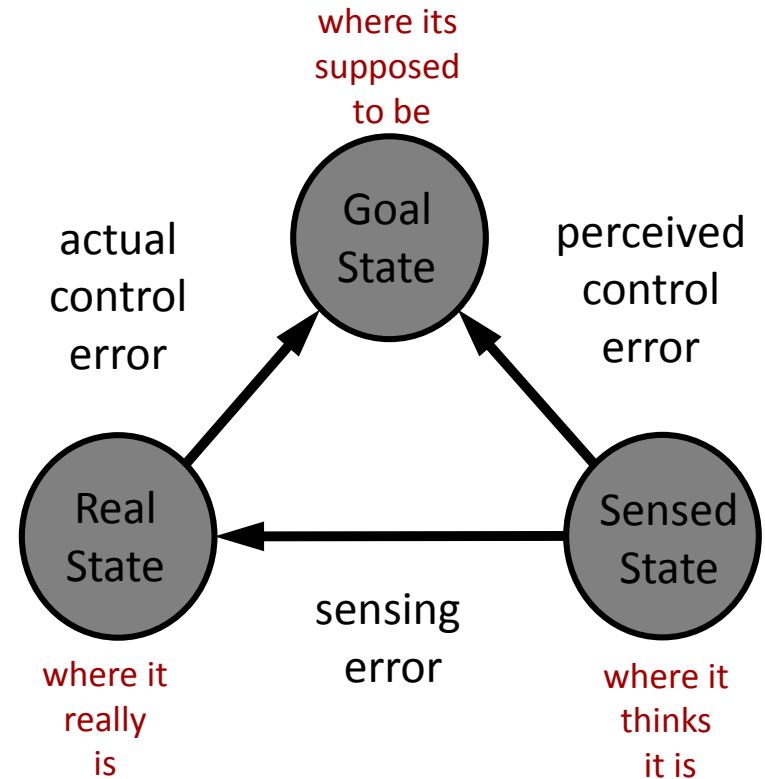
- Derivatives can blow up:
 - When the system is moved while not under control.
- Derivatives can also attenuate:
 - When system is stopped for long period of time.





Control Versus Sensor Error

- The robot cannot tell if its pose is in error.
- **Sensing Error =**
real pose – sensed pose
- **Perceived Control Error =**
goal pose – sensed pose.
- **Actual Control Error =**
goal pose – real pose.
- Your job is usually to make sure the perceived control error is small.
- How can you make sensing error smaller?





Agenda Today

CarnegieMellon
THE ROBOTICS INSTITUTE

- Administrative Issues
- Feedback Control
- Tuning
- **Model Reference Control**
- Identification
- Feedforward Control
- 2 Dof Control
- Overview of Lab 4



Weal World Wobots (are waskewy)




- Effect 1: Dynamics
 - It often takes much longer to do anything because real systems have dynamics (are low pass filters).
- Effect 2: Comms Delays:
 - It takes two cycles for response to commands generated in this cycle to be measurable.
 - **What is a simple way to compensate for delays?**
- Effect 3: Hierarchy / Cascading
 - There are lower control levels below yours. See next.



Effect 1: Dynamics (Momentum)

Carnegie Mellon
THE ROBOTICS INSTITUTE



- Robots have mass and hence momentum.
 - A form of filter you have little control over.
- They **do not follow their (velocity, position) “commands” precisely.**
 - Yet, Newton’s laws will certainly hold.
 - If you **knew** the **actual** forces and inertias and frictions ...
- Its both good and bad
 - Attenuates disturbances
 - Attenuates your commands





Effect 2: Delays


- Do not expect the feedback to change for a few cycles after a new command is issued.





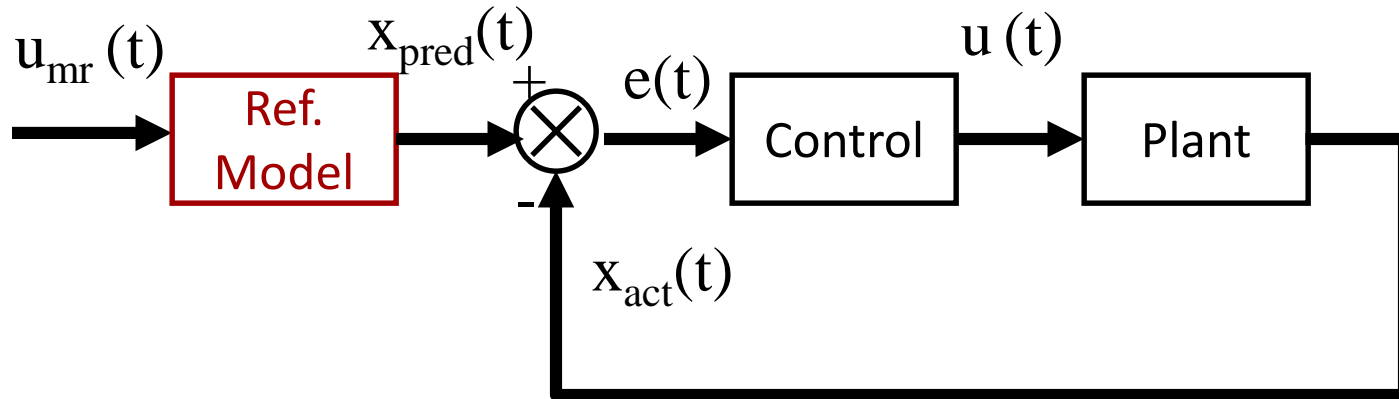
Effect 3: Lower Control Levels

- Commanded acceleration is limited.
 - Response for a step speed input is often a ramp.
- Commanded speed is limited.
- That gives an odd sort of filter for a continuously varying input.
- Actual robot response to the black line is a separate matter.





Model Reference Control



$$e(t) = x_{pred}(t) - x_{act}(t)$$



$$u(t) = k_p e(t) + k_d \frac{d}{dt} e(t) + k_i \int e(t) dt$$

- Basic idea is to form errors w.r.t. **predicted** performance of “reference model”.



Example

- Model Reference Control of Point Mass.
 - For a little better model



- 2.5 secs. versus 5 secs. (for critically damped feedback) and no overshoot!
 - Maybe this is a good idea!
 - Where does the reference model come from?
 - 1: Make it up
 - 2: Model the real system – system identification.



Agenda Today

CarnegieMellon
THE ROBOTICS INSTITUTE

- Administrative Issues
- Feedback Control
- Tuning
- Model Reference Control
- **Identification**
- Feedforward Control
- 2 Dof Control
- Overview of Lab 4




Rotary Identification

(Example From Nomad Robot)

Carnegie Mellon
THE ROBOTICS INSTITUTE

- Notice:
 - Delay (~ 0.15 secs)
 - Acceleration limit ($= \sim 3$ rads per sec per sec)






Linear Identification

(Example From Nomad Robot)

Carnegie Mellon
THE ROBOTICS INSTITUTE

- Notice:
 - Delay (~ 0.15 secs)
 - Acceleration limit ($= 0.5$ m.sec)





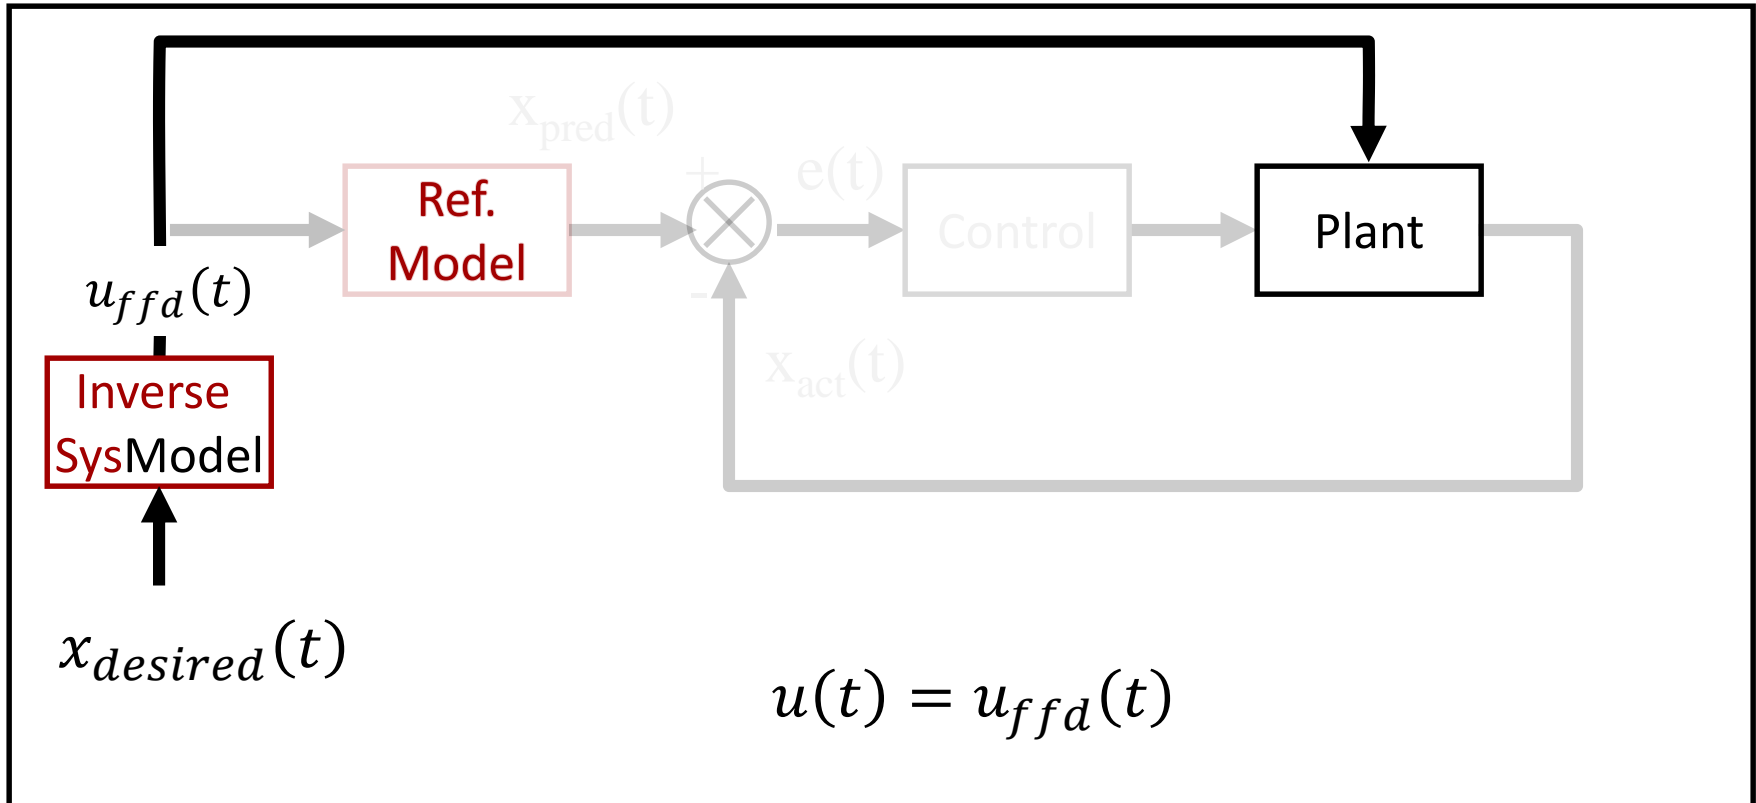
Agenda Today

CarnegieMellon
THE ROBOTICS INSTITUTE

- Administrative Issues
- Feedback Control
- Tuning
- Model Reference Control
- Identification
- **Feedforward Control**
- 2 Dof Control
- Overview of Lab 4



Feedforward Control





- Basic idea is know **what to ask for in order to get what you want.**
 - Not only do you have a reference model, **but you invert it too.**
- What is $x_{desired}(t)$ you ask? – any curve that starts where you are and ends where you want.



Inverting is not so easy....



- How do you invert a differential equation?
- I.e. how do you choose the input that produces you output you want?
- There **may be no solution**....





Example – that really works

- Feedforward Control of Point Mass
 - In this case it is also a “bang bang” control



- This is the time optimal control (**2** secs) for this case.
- Issue: What if there is (unknown) friction?



Agenda Today

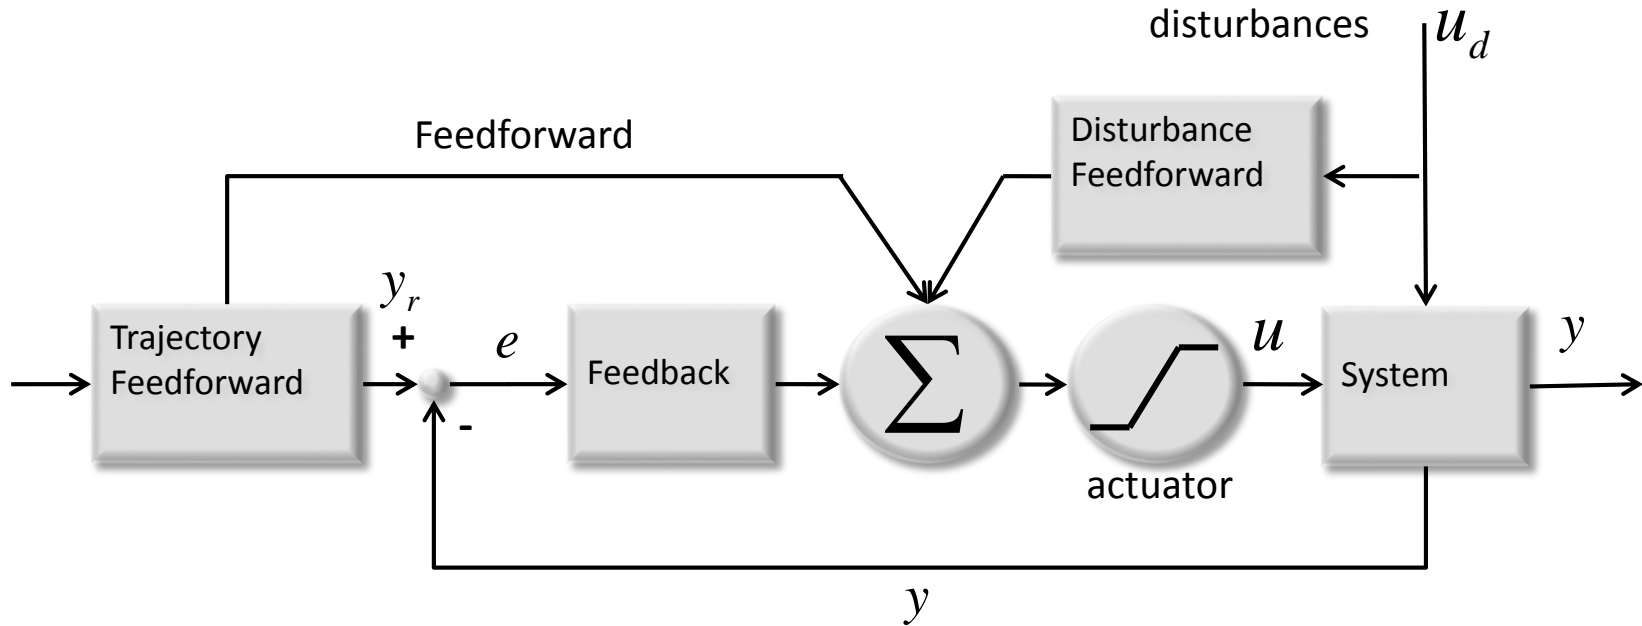
CarnegieMellon
THE ROBOTICS INSTITUTE

- Administrative Issues
- Feedback Control
- Tuning
- Model Reference Control
- Identification
- Feedforward Control
- **2 Dof Control**
- Overview of Lab 4



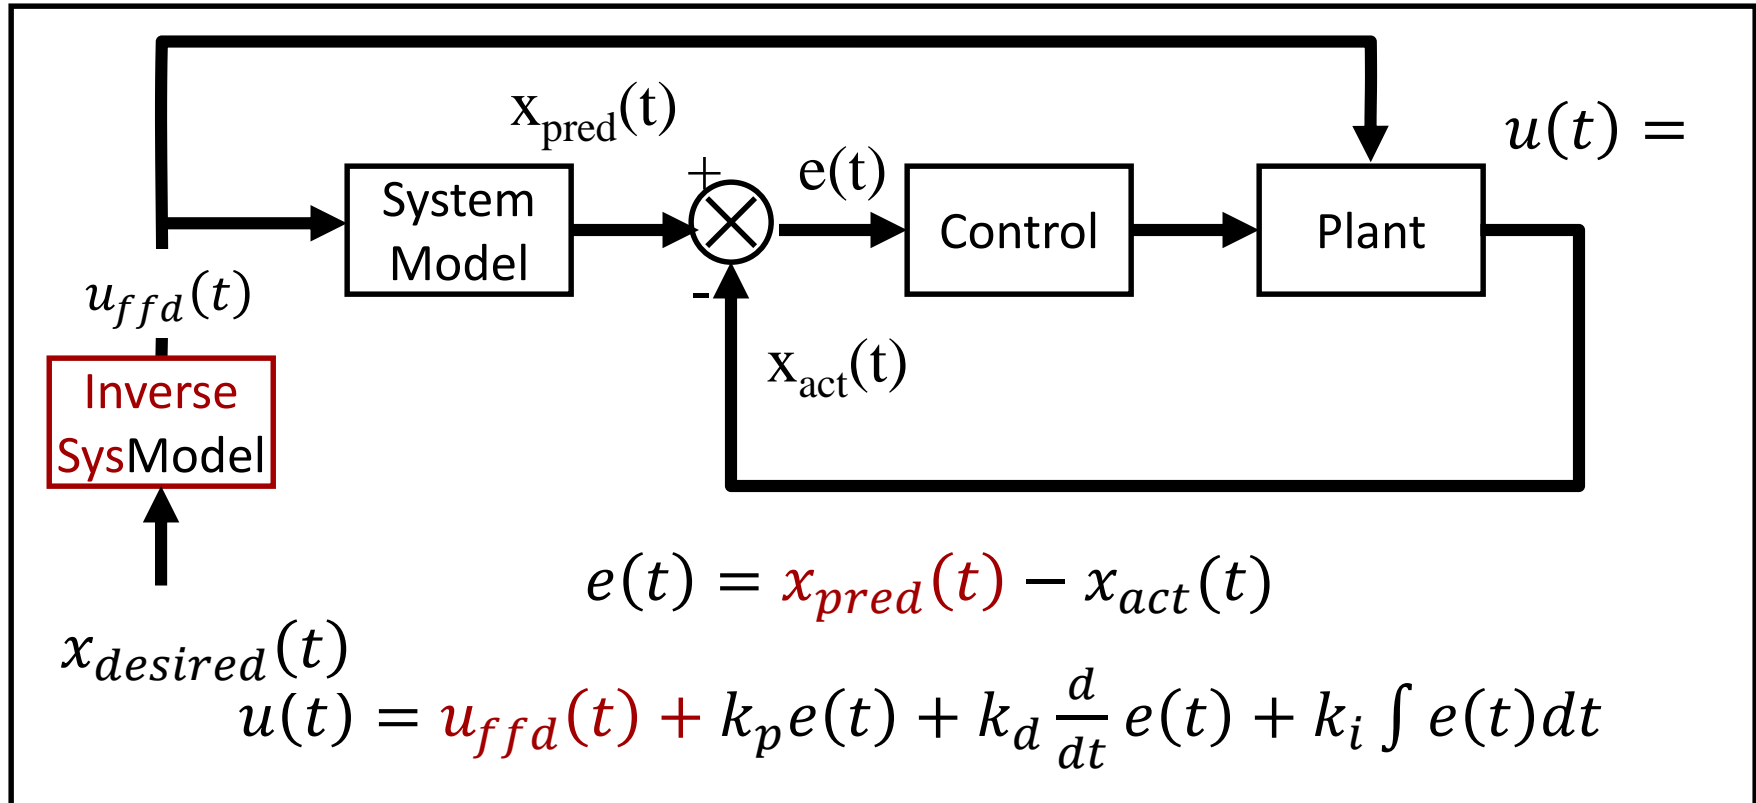
Two Degree of Freedom (Feedforward with Feedback Trim)

- Do your best to determine the input which causes the system to go where you want open loop.
 - Predict even the disturbances if that is possible.
- Form errors on the **predicted** response.





2 DOF Control




- Controller has a feedforward term **and** a feedback term.
 - Feedforward works perfectly if there are no errors.
 - Feedback removes the (small) errors if any occur.



Final Result

- An optimal trajectory that **also** rejects disturbances.






Bottom Line

	Feedback	Feedforward
Removes Unpredictable Errors and Disturbances	(+) YES	(-) NO
Removes Predictable Errors and Disturbances	(-) NO	(+) YES
Removes Errors and Disturbances Before They Happen	(-) NO	(+) YES
Requires Model of System	(+) NO	(-) YES
Affects Stability of System	(-) YES	(+) NO



This Lab

- Command a fwd trapezoidal velocity profile:
 - with a terminal $V=0$ period 1 second long.
- Assume the system will follow a trapezoidal velocity profile with delay.
- Compare the real response to the “expected” and remove the errors with feedback.





Preparation

Carnegie Mellon
THE ROBOTICS INSTITUTE

- Did everyone read lab 4 before today?



Lab 4 Preparation

Carnegie Mellon
THE ROBOTICS INSTITUTE

- [Click for Lab4 Writeup](#)